



Leveling Up My Linux Kernel Contributions

- Troubleshooting the kernel panic

Juhee Kang



Speaker



Juhee Kang (claudiajkang@gmail.com)

Open-Source Developer

Project: Linux Kernel Networking Stack, Kubernetes i18n

Interest: Kernel, Cloud, Network, Backend, etc...

Special Lecture at Universities

- Open-Source Software Development Mechanism.



Sometimes kernel panic occurs

```
QEMU
[ 4.251820] 0b00 1048575 sr0
[ 4.251859] driver: sr
[ 4.252046] Kernel panic - not syncing: VFS: Unable to mount root fs on unknown-block(0,0)
[ 4.252391] CPU: 0 PID: 1 Comm: swapper/0 Not tainted 4.12.3 #1
[ 4.302283] Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS rel-1.11.1-0-g0551a4be2c-prebuilt.qemu-project.org 04/01/2014
[ 4.302631] Call Trace:
[ 4.303823] dump_stack+0x1d/0x70
[ 4.303991] panic+0xca/0x203
[ 4.304154] ? printk+0x3e/0x46
[ 4.304305] mount_block_root+0x171/0x221
[ 4.304401] ? set_debug_rodatta+0x12/0x12
[ 4.304496] mount_root+0x101/0x10a
[ 4.304590] ? set_debug_rodatta+0x12/0x12
[ 4.304639] prepare_namespace+0x134/0x16c
[ 4.304685] kernel_init_freeable+0x1da/0x1f3
[ 4.304730] ? rest_init+0x80/0x80
[ 4.304775] kernel_init+0x9/0x100
[ 4.304818] ret_from_fork+0x22/0x30
[ 4.305303] Kernel Offset: 0x37200000 from 0xffffffff81000000 (relocation range: 0xffffffff80000000-0xffffffffbfffffff)
[ 4.305632] ---[ end Kernel panic - not syncing: VFS: Unable to mount root fs on unknown-block(0,0) ]---
```



```
[ 79.803379] [T969] sysrq: Trigger a crash
[ 79.803924] [T969] Kernel panic - not syncing: sysrq triggered crash
[ 79.804578] [T969] CPU: 6 PID: 969 Comm: bash Not tainted 6.2.0-rc2 #2 e5c88a38d28ce676364b6ebf0...
[ 79.805639] [T969] Hardware name: QEMU Standard PC (Q35 + ICH9, 2009), BIOS 1.15.0-1 04/01/2014
[ 79.806640] [T969] Call Trace:
[ 79.806950] [T969] <TASK>
[ 79.807229] [T969] dump_stack_lvl+0x49/0x61
[ 79.807693] [T969] panic+0x113/0x28c
[ 79.808072] [T969] sysrq_handle_crash+0x15/0x20
[ 79.808583] [T969] __handle_sysrq.cold+0x51/0x13c
[ 79.809205] [T969] write_sysrq_trigger+0x42/0x50
[ 79.809668] [T969] proc_reg_write+0x54/0xa0
[ 79.810195] [T969] ? rcu_read_lock_any_held+0x79/0xa0
[ 79.810750] [T969] vfs_write+0xc7/0x4a0
[ 79.811157] [T969] ? __do_sys_newfstatat+0x35/0x60
[ 79.811672] [T969] ? lock_is_held_type+0xe1/0x140
[ 79.812167] [T969] ksys_write+0x67/0xf0
[ 79.812575] [T969] do_syscall_64+0x3f/0x90
[ 79.813010] [T969] entry_SYSCALL_64_after_hwframe+0x72/0xdc
[ 79.813581] [T969] RIP: 0033:0x7f6867fcea37
[ 79.814029] [T969] Code: 10 00 f7 d8 64 89 02 48 c7 c0 ff ff ff ff eb b7 0f 1f 00 f3 0f 1e fa 64 8b 0
4 25 18 00 00 00 85 c0 75 10 b8 01 00 00 00 0f 05 <48> 3d 00 f0 ff ff 77 51 c3 48 83 ec 28 48 89 54 ...
[ 79.816012] [T969] RSP: 002b:00007ffcd75d2a18 EFLAGS: 00000246 ORIG_RAX: 0000000000000001
[ 79.816832] [T969] RAX: ffffffff80000000 RBX: 0000000000000002 RCX: 00007f6867fcea37
[ 79.817605] [T969] RDX: 0000000000000002 RSI: 000055cccc98b330 RDI: 0000000000000001
[ 79.818378] [T969] RBP: 000055cccc98b330 R08: 0000000000000000 R09: 000055cccc98b330
[ 79.819154] [T969] R10: 00007f68680d3d60 R11: 0000000000000246 R12: 0000000000000002
[ 79.819935] [T969] R13: 00007f68680d4780 R14: 00007f68680d0600 R15: 00007f68680cfa00
[ 79.820728] [T969] </TASK>
[ 79.821439] [T969] Kernel Offset: disabled
[ 79.821895] [T969] ---[ end Kernel panic - not syncing: sysrq triggered crash ]---
```




Panic





Normally I solved the kernel panic



Infinite loop...



Troubleshooting kernel panic?



Is there any effective way of troubleshooting kernel panic?



Session's goal

Step through **kernel panic debugging** with example!



How to cause kernel panic with a single command?

Easiest way to raise **kernel panic**

```
# echo c > /proc/sysrq-trigger
```

<https://www.kernel.org/doc/html/v6.2-rc2/admin-guide/sysrq.html>



Raise Kernel panic

```
# echo c > /proc/sysrq-trigger
```

```
[ 79.803379][ T969] sysrq: Trigger a crash
[ 79.803924][ T969] Kernel panic - not syncing: sysrq triggered crash
[ 79.804578][ T969] CPU: 6 PID: 969 Comm: bash Not tainted 6.2.0-rc2 #2 e5c88a38d28ce676364b6ebf0...
[ 79.805639][ T969] Hardware name: QEMU Standard PC (Q35 + ICH9, 2009), BIOS 1.15.0-1 04/01/2014
[ 79.806640][ T969] Call Trace:
[ 79.806950][ T969] <TASK>
[ 79.807229][ T969] dump_stack_lvl+0x49/0x61
[ 79.807693][ T969] panic+0x113/0x28c
[ 79.808072][ T969] sysrq_handle_crash+0x15/0x20
[ 79.808583][ T969] __handle_sysrq.cold+0x51/0x13c
[ 79.809205][ T969] write_sysrq_trigger+0x42/0x50
[ 79.809668][ T969] proc_reg_write+0x54/0xa0
[ 79.810195][ T969] ? rcu_read_lock_any_held+0x79/0xa0
[ 79.810750][ T969] vfs_write+0xc7/0x4a0
[ 79.811157][ T969] ? __do_sys_newfstatat+0x35/0x60
[ 79.811672][ T969] ? lock_is_held_type+0xe1/0x140
[ 79.812167][ T969] ksys_write+0x67/0xf0
[ 79.812575][ T969] do_syscall_64+0x3f/0x90
[ 79.813010][ T969] entry_SYSCALL_64_after_hwframe+0x72/0xdc
[ 79.813581][ T969] RIP: 0033:0x7f6867fcea37
[ 79.814029][ T969] Code: 10 00 f7 d8 64 89 02 48 c7 c0 ff ff ff ff eb b7 0f 1f 00 f3 0f 1e fa 64 8b 0
4 25 18 00 00 00 85 c0 75 10 b8 01 00 00 00 0f 05 <48> 3d 00 f0 ff ff 77 51 c3 48 83 ec 28 48 89 54 ...
[ 79.816012][ T969] RSP: 002b:00007ffcd75d2a18 EFLAGS: 00000246 ORIG_RAX: 0000000000000001
[ 79.816832][ T969] RAX: ffffffffda RBX: 0000000000000002 RCX: 00007f6867fcea37
[ 79.817605][ T969] RDX: 0000000000000002 RSI: 000055cccc98b330 RDI: 0000000000000001
[ 79.818378][ T969] RBP: 000055cccc98b330 R08: 0000000000000000 R09: 000055cccc98b330
[ 79.819154][ T969] R10: 00007f68680d3d60 R11: 0000000000000246 R12: 0000000000000002
[ 79.819935][ T969] R13: 00007f68680d4780 R14: 00007f68680d0600 R15: 00007f68680cfa00
[ 79.820728][ T969] </TASK>
[ 79.821439][ T969] Kernel Offset: disabled
[ 79.821895][ T969] ---[ end Kernel panic - not syncing: sysrq triggered crash ]---
```



```
# echo c > /proc/sysrq-trigger
```

But Why?



Analysis kernel panic log

First, let's look at the Kernel Panic log



Analysis kernel panic log

① Skim through Kernel Panic log

```
[ 79.803379][ T969] sysrq: Trigger a crash
[ 79.803924][ T969] Kernel panic - not syncing: sysrq triggered crash
[ 79.804578][ T969] CPU: 6 PID: 969 Comm: bash Not tainted 6.2.0-rc2 #2 e5c88a38d28ce676364b6ebf0...
[ 79.805639][ T969] Hardware name: QEMU Standard PC (Q35 + ICH9, 2009), BIOS 1.15.0-1 04/01/2014
[ 79.806640][ T969] Call Trace:
[ 79.806950][ T969] <TASK>
[ 79.807229][ T969] dump_stack_lvl+0x49/0x61
[ 79.807693][ T969] panic+0x113/0x28c
[ 79.808072][ T969] sysrq_handle_crash+0x15/0x20
[ 79.808583][ T969] __handle_sysrq.cold+0x51/0x13c
[ 79.809205][ T969] write_sysrq_trigger+0x42/0x50
[ 79.809668][ T969] proc_reg_write+0x54/0xa0
[ 79.810195][ T969] ? rcu_read_lock_any_held+0x79/0xa0
[ 79.810750][ T969] vfs_write+0xc7/0x4a0
[ 79.811157][ T969] ? __do_sys_newfstatat+0x35/0x60
[ 79.811672][ T969] ? lock_is_held_type+0xe1/0x140
[ 79.812167][ T969] ksys_write+0x67/0xf0
[ 79.812575][ T969] do_syscall_64+0x3f/0x90
[ 79.813010][ T969] entry_SYSCALL_64_after_hwframe+0x72/0xdc
[ 79.813581][ T969] RIP: 0033:0x7f6867fcea37
[ 79.814029][ T969] Code: 10 00 f7 d8 64 89 02 48 c7 c0 ff ff ff ff eb b7 0f 1f 00 f3 0f 1e fa 64 8b 0
4 25 18 00 00 00 85 c0 75 10 b8 01 00 00 00 0f 05 <48> 3d 00 f0 ff ff 77 51 c3 48 83 ec 28 48 89 54 ...
[ 79.816012][ T969] RSP: 002b:00007ffcd75d2a18 EFLAGS: 00000246 ORIG_RAX: 0000000000000001
[ 79.816832][ T969] RAX: ffffffffda RBX: 0000000000000002 RCX: 00007f6867fcea37
[ 79.817605][ T969] RDX: 0000000000000002 RSI: 000055cccc98b330 RDI: 0000000000000001
[ 79.818378][ T969] RBP: 000055cccc98b330 R08: 0000000000000000 R09: 000055cccc98b330
[ 79.819154][ T969] R10: 00007f68680d3d60 R11: 0000000000000246 R12: 0000000000000002
[ 79.819935][ T969] R13: 00007f68680d4780 R14: 00007f68680d0600 R15: 00007f68680cfa00
[ 79.820728][ T969] </TASK>
[ 79.821439][ T969] Kernel Offset: disabled
[ 79.821895][ T969] ---[ end Kernel panic - not syncing: sysrq triggered crash ]---
```



Analysis kernel panic log

① Skim through Kernel Panic log

Panic log header provides the abstraction about the crash

ex) Cause of the panic

```
[ 79.803379][ T969] sysrq: Trigger a crash
[ 79.803924][ T969] Kernel panic - not syncing: sysrq triggered crash
[ 79.804578][ T969] CPU: 6 PID: 969 Comm: bash Not tainted 6.2.0-rc2 #2 e5c88a38d28ce676364b6ebf0...
[ 79.805639][ T969] Hardware name: QEMU Standard PC (Q35 + ICH9, 2009), BIOS 1.15.0-1 04/01/2014
[ 79.806640][ T969] Call trace:
[ 79.806950][ T969] <TASK>
[ 79.807229][ T969] dump_stack_lvl+0
[ 79.807693][ T969] panic+0x113/0x28c
[ 79.808072][ T969] sysrq_handle_crash+0x15/0x20
[ 79.808583][ T969] __handle_sysrq.cold+0x51/0x13c
[ 79.809205][ T969] write_sysrq_trigger+0x42/0x50
[ 79.809668][ T969] proc_reg_write+0x54/0xa0
[ 79.810195][ T969] ? rcu_read_lock_any_held+0x79/0xa0
[ 79.810750][ T969] vfs_write+0xc7/0x4a0
[ 79.811157][ T969] ? __do_sys_newfstatat+0x35/0x60
[ 79.811672][ T969] ? lock_is_held_type+0xe1/0x140
[ 79.812167][ T969] ksys_write+0x67/0xf0
[ 79.812575][ T969] do_syscall_64+0x3f/0x90
[ 79.813010][ T969] entry_SYSCALL_64_after_hwframe+0x72/0xdc
[ 79.813581][ T969] RIP: 0033:0x7f6867fcea37
[ 79.814029][ T969] Code: 10 00 f7 d8 64 89 02 48 c7 c0 ff ff ff ff eb b7 0f 1f 00 f3 0f 1e fa 64 8b 0
4 25 18 00 00 00 85 c0 75 10 b8 01 00 00 00 0f 05 <48> 3d 00 f0 ff ff 77 51 c3 48 83 ec 28 48 89 54 ...
[ 79.816012][ T969] RSP: 002b:00007ffcd75d2a18 EFLAGS: 00000246 ORIG_RAX: 0000000000000001
[ 79.816832][ T969] RAX: ffffffffda RBX: 0000000000000002 RCX: 00007f6867fcea37
[ 79.817605][ T969] RDX: 0000000000000002 RSI: 000055cccc98b330 RDI: 0000000000000001
[ 79.818378][ T969] RBP: 000055cccc98b330 R08: 0000000000000000 R09: 000055cccc98b330
[ 79.819154][ T969] R10: 00007f68680d3d60 R11: 0000000000000246 R12: 0000000000000002
[ 79.819935][ T969] R13: 00007f68680d4780 R14: 00007f68680d0600 R15: 00007f68680cfa00
[ 79.820728][ T969] </TASK>
[ 79.821439][ T969] Kernel Offset: disabled
[ 79.821895][ T969] ---[ end Kernel panic - not syncing: sysrq triggered crash ]---
```

Panic log header



Analysis kernel panic log

① Skim through Kernel Panic log

Panic log header provides the abstraction about the crash

Call Trace provides the context information about the crash

dump_stack_lvl + 0x49 / 0x61



Function Symbol

Offset to the code

Code size of the func.

```
[ 79.803379][ T969] sysrq: Trigger a crash
[ 79.803924][ T969] Kernel panic - not syncing: sysrq triggered crash
[ 79.804578][ T969] CPU: 6 PID: 969 Comm: bash Not tainted 6.2.0-rc2 #2 e5c88a38d28ce676364b6ebf0...
[ 79.805639][ T969] Hardware name: QEMU Standard PC (Q35 + ICH9, 2009), BIOS 1.15.0-1 04/01/2014
[ 79.806640][ T969]
[ 79.806950][ T969] Call Trace:
[ 79.807229][ T969] <TASK>
[ 79.807229][ T969] dump_stack_lvl+0x49/0x61
[ 79.807693][ T969] panic+0x113/0x28c
[ 79.808072][ T969] sysrq_handle_crash+0x15/0x20
[ 79.808583][ T969] __handle_sysrq.cold+0x51/0x13c
[ 79.809205][ T969] write_sysrq_trigger+0x42/0x50
[ 79.809668][ T969] proc_reg_write+0x54/0xa0
[ 79.810195][ T969] ? rcu_read_lock_any_held+0x79/0xa0
[ 79.810750][ T969] vfs_write+0xc7/0x4a0
[ 79.811157][ T969] ? __do_sys_newfstatat+0x35/0x60
[ 79.811672][ T969] ? lock_is_held_type+0xe1/0x140
[ 79.812167][ T969] ksys_write+0x67/0xf0
[ 79.812575][ T969] do_syscall_64+0x3f/0x90
[ 79.813010][ T969] entry_SYSCALL_64_after_hwframe+0x72/0xdc
[ 79.813581][ T969] RIP: 0033:0x7f6867fcea57
[ 79.814029][ T969] Code: 10 00 f7 d8 64 89 02 48 c7 c0 ff ff ff ff eb b7 0f 1f 00 f3 0f 1e fa 64 8b 0
4 25 18 00 00 00 85 c0 75 10 b8 01 00 00 00 0f 05 <48> 3d 00 f0 ff ff 77 51 c3 48 83 ec 28 48 89 54 ...
[ 79.816012][ T969] RSP: 002b:00007ffcd75d2a18 EFLAGS: 00000246 ORIG_RAX: 0000000000000001
[ 79.816832][ T969] RAX: ffffffffda RBX: 0000000000000002 RCX: 00007f6867fcea37
[ 79.817605][ T969] RDX: 0000000000000002 RSI: 000055cccc98b330 RDI: 0000000000000001
[ 79.818378][ T969] RBP: 000055cccc98b330 R08: 0000000000000000 R09: 000055cccc98b330
[ 79.819154][ T969] R10: 00007f68680d3d60 R11: 0000000000000246 R12: 0000000000000002
[ 79.819935][ T969] R13: 00007f68680d4780 R14: 00007f68680d0600 R15: 00007f68680cfa00
[ 79.820728][ T969] </TASK>
[ 79.821439][ T969] Kernel Offset: disabled
[ 79.821895][ T969] ---[ end Kernel panic - not syncing: sysrq triggered crash ]---
```

Call Trace



Analysis kernel panic log

① Skim through Kernel Panic log

Panic log header provides the abstraction about the crash

Call Trace provides the context information about the crash

Register Information provides the current executing dump of CPU registers

```
[ 79.803379][ T969] sysrq: Trigger a crash
[ 79.803924][ T969] Kernel panic - not syncing: sysrq triggered crash
[ 79.804578][ T969] CPU: 6 PID: 969 Comm: bash Not tainted 6.2.0-rc2 #2 e5c88a38d28ce676364b6ebf0...
[ 79.805639][ T969] Hardware name: QEMU Standard PC (Q35 + ICH9, 2009), BIOS 1.15.0-1 04/01/2014
[ 79.806640][ T969] Call Trace:
[ 79.806950][ T969] <TASK>
[ 79.807229][ T969] dump_stack_lvl+0x49/0x61
[ 79.807693][ T969] panic+0x113/0x28c
[ 79.808072][ T969] sysrq_handle_crash+0x15/0x20
[ 79.808583][ T969] __handle_sysrq.cold+0x51/0x13c
[ 79.809205][ T969] write_sysrq_trigger+0x42/0x50
[ 79.809668][ T969] proc_reg_write+0x54/0xa0
[ 79.810195][ T969] ? rcu_read_lock_any_held+0x79/0xa0
[ 79.810750][ T969] vfs_write+0xc7/0x4a0
[ 79.811157][ T969] ? __do_sys_newfstatat+0x35/0x60
[ 79.811672][ T969] ? lock_is_held_type+0xe1/0x140
[ 79.812167][ T969] ksys_write+0x67/0xf0
[ 79.812575][ T969] do_syscall_64+0x3f/0x90
[ 79.813010][ T969] entry_SYSCALL_64_after_hwframe+0x72/0x7d
[ 79.813581][ T969] RIP: 0033:0x7f6867fcea37
[ 79.814029][ T969] Code: 10 00 f7 d8 64 89 02 48 c7 c0 ff ff ff ff eb b7 0f 1f 00 f3 0f 1e fa 64 8b 0
4 25 18 00 00 00 85 c0 75 10 b8 01 00 00 00 0f 05 <48> 3d 00 f0 ff ff 77 51 c3 48 83 ec 28 48 89 54 ...
[ 79.816012][ T969] RSP: 002b:00007ffcd75d2a18 EFLAGS: 00000246 ORIG_RAX: 0000000000000001
[ 79.816832][ T969] RAX: ffffffffda RBX: 0000000000000002 RCX: 00007f6867fcea37
[ 79.817605][ T969] RDX: 0000000000000002 RSI: 000055cccc98b330 RDI: 0000000000000001
[ 79.818378][ T969] RBP: 000055cccc98b330 R08: 0000000000000000 R09: 000055cccc98b330
[ 79.819154][ T969] R10: 00007f68680d3d60 R11: 0000000000000246 R12: 0000000000000002
[ 79.819935][ T969] R13: 00007f68680d4780 R14: 00007f68680d0600 R15: 00007f68680cfa00
[ 79.820728][ T969] </TASK>
[ 79.821439][ T969] Kernel Offset: disabled
[ 79.821895][ T969] ---[ end Kernel panic - not syncing: sysrq triggered crash ]---
```

Register Info

```
RIP: 0033:0x7f6867fcea37
Code: 10 00 f7 d8 64 89 02 48 c7 c0 ff ff ff ff eb b7 0f 1f 00 f3 0f 1e fa 64 8b 0
4 25 18 00 00 00 85 c0 75 10 b8 01 00 00 00 0f 05 <48> 3d 00 f0 ff ff 77 51 c3 48 83 ec 28 48 89 54 ...
RSP: 002b:00007ffcd75d2a18 EFLAGS: 00000246 ORIG_RAX: 0000000000000001
RAX: ffffffffda RBX: 0000000000000002 RCX: 00007f6867fcea37
RDX: 0000000000000002 RSI: 000055cccc98b330 RDI: 0000000000000001
RBP: 000055cccc98b330 R08: 0000000000000000 R09: 000055cccc98b330
R10: 00007f68680d3d60 R11: 0000000000000246 R12: 0000000000000002
R13: 00007f68680d4780 R14: 00007f68680d0600 R15: 00007f68680cfa00
```




Analysis kernel panic log

① Skim through Kernel Panic log

Panic log header provides the abstraction about the crash

Call Trace provides the context information about the crash

Register Information provides the current executing dump of CPU registers

RIP holds the current executing instruction

Code includes the current executing code information

```
[ 79.803379][ T969] sysrq: Trigger a crash
[ 79.803924][ T969] Kernel panic - not syncing: sysrq triggered crash
[ 79.804578][ T969] CPU: 6 PID: 969 Comm: bash Not tainted 6.2.0-rc2 #2 e5c88a38d28ce676364b6ebf0...
[ 79.805639][ T969] Hardware name: QEMU Standard PC (Q35 + ICH9, 2009), BIOS 1.15.0-1 04/01/2014
[ 79.806640][ T969] Call Trace:
[ 79.806950][ T969] <TASK>
[ 79.807229][ T969] dump_stack_lvl+0x49/0x61
[ 79.807693][ T969] panic+0x113/0x28c
[ 79.808072][ T969] sysrq_handle_crash+0x15/0x20
[ 79.808583][ T969] __handle_sysrq.cold+0x51/0x13c
[ 79.809205][ T969] write_sysrq_trigger+0x42/0x50
[ 79.809668][ T969] proc_reg_write+0x54/0xa0
[ 79.810195][ T969] ? rcu_read_lock_any_held+0x79/0xa0
[ 79.810750][ T969] vfs_write+0xc7/0x4a0
[ 79.811157][ T969] ? __do_sys_newfstatat+0x35/0x60
[ 79.811672][ T969] ? lock_is_held_type+0xe1/0x140
[ 79.812167][ T969] ksys_write+0x67/0xf0
[ 79.812575][ T969] do_syscall_64+0x3f/0x90
[ 79.813010][ T969] entry_SYSCALL_64_after_hwframe+0x72/0xdc
[ 79.813581][ T969]
[ 79.814029][ T969] RIP: 0033:0x7f6867fcea37
[ 79.814029][ T969] Code: 10 00 f7 d8 64 89 02 48 c7 c0 ff ff ff ff eb b7 0f 1f 00 f3 0f 1e fa 64 8b 0
[ 79.814029][ T969] c0 75 10 b8 01 00 00 00 0f 05 <48> 3d 00 f0 ff ff 77 51 c3 48 83 ec 28 48 89 54 ...
[ 79.816012][ T969] RSP: 002b:00007ffcd75d2a18 EFLAGS: 00000246 ORIG_RAX: 0000000000000001
[ 79.816832][ T969] RAX: ffffffffda RBX: 0000000000000002 RCX: 00007f6867fcea37
[ 79.817605][ T969] RDX: 0000000000000002 RSI: 000055cccc98b330 RDI: 0000000000000001
[ 79.818378][ T969] RBP: 000055cccc98b330 R08: 0000000000000000 R09: 000055cccc98b330
[ 79.819154][ T969] R10: 00007f68680d3d60 R11: 0000000000000246 R12: 0000000000000002
[ 79.819935][ T969] R13: 00007f68680d4780 R14: 00007f68680d0600 R15: 00007f68680cfa00
[ 79.820728][ T969] </TASK>
[ 79.821439][ T969] Kernel Offset: disabled
[ 79.821895][ T969] ---[ end Kernel panic - not syncing: sysrq triggered crash ]---
```

Register Info



For call trace analysis, save the panic log



② Decode call trace

```
# ./scripts/decode_stacktrace.sh vmlinux < err.log
```

Script from kernel source Current Kernel Image Crash Log

```
[79.803379][ T969] sysrq: Trigger a crash  
[79.803924][ T969] Kernel panic - not syncing: sysrq triggered crash  
[79.804578][ T969] CPU: 6 PID: 969 Comm: bash Not tainted 6.2.0-rc2 #2 e5c88a38d28ce676364b6ebf0...  
[79.805639][ T969] Hardware name: QEMU Standard PC (Q35 + ICH9, 2009), BIOS 1.15.0-1 04/01/2014  
[79.806640][ T969] Call Trace:  
[79.806950][ T969] <TASK>  
[79.807229][ T969] dump_stack_lvl+0x49/0x61  
[79.807693][ T969] panic+0x113/0x28c  
[79.808072][ T969] sysrq_handle_crash+0x15/0x20  
[79.808583][ T969] __handle_sysrq.cold+0x51/0x13c  
[79.809205][ T969] write_sysrq_trigger  
[79.809668][ T969] proc_reg_write+0x...  
[79.810195][ T969] ? rcu_read_lock_any  
[79.810750][ T969] vfs_write+0xc7/0x4...  
[79.811157][ T969] ? __do_sys_newfstab  
[79.811672][ T969] ? lock_is_held_type  
[79.812167][ T969] ksys_write+0x67/0x...  
[79.812575][ T969] do_syscall_64+0x3f...  
[79.813010][ T969] entry_SYSCALL_64...  
[79.813581][ T969] RIP: 0033:0x7f6867f...  
[79.814029][ T969] Code: 10 00 f7 d8 64 ... b7 0f 1f 00 f3 0f 1e fa 64 8b 0  
4 25 18 00 00 00 85 c0 75 10 b8 01 00 00 ... 77 51 c3 48 83 ec 28 48 89 54 ...  
[79.816012][ T969] RSP: 002b:00007ffcd75d2a18 EFLAGS: 00000246 ORIG_RAX: 0000000000000001  
[79.816832][ T969] RAX: ffffffffda RBX: 0000000000000002 RCX: 00007f6867fcea37  
[79.817605][ T969] RDY: 0000000000000002 RSI: 000055cccc98b330 RDI: 0000000000000001  
[79.818378][ T969] RBP: 000055cccc98b330 R08: 0000000000000000 R09: 000055cccc98b330  
[79.819154][ T969] R10: 00007f68680d3d60 R11: 0000000000000246 R12: 0000000000000002  
[79.819935][ T969] R13: 00007f68680d4780 R14: 00007f68680d0600 R15: 00007f68680cfa00  
[79.820728][ T969] </TASK>  
[79.821439][ T969] Kernel Offset: disabled  
[79.821895][ T969] ---[ end Kernel panic - not syncing: sysrq triggered crash ]---
```



Analysis kernel panic log

① Skim through Kernel Panic log

② Decode call trace

```
[ 79.806640][ T969] Call Trace:
[ 79.806950][ T969] <TASK>
[ 79.807229][ T969] dump_stack_lvl+0x49/0x61
[ 79.807693][ T969] panic+0x113/0x28c
[ 79.808072][ T969] sysrq_handle_crash+0x15/0x20
[ 79.808583][ T969] __handle_sysrq.cold+0x51/0x13c
[ 79.809205][ T969] write_sysrq_trigger+0x42/0x50
[ 79.809668][ T969] proc_reg_write+0x54/0xa0
[ 79.810195][ T969] ? rcu_read_lock_any_held+0x79/0xa0
[ 79.810750][ T969] vfs_write+0xc7/0x4a0
[ 79.811157][ T969] ? __do_sys_newfstatat+0x35/0x60
[ 79.811672][ T969] ? lock_is_held_type+0xe1/0x140
[ 79.812167][ T969] ksys_write+0x67/0xf0
[ 79.812575][ T969] do_syscall_64+0x3f/0x90
[ 79.813010][ T969] entry_SYSCALL_64_after_hwframe+0x72/0xdc
```



```
[ 79.803379][ T969] sysrq: Trigger a crash
[ 79.803924][ T969] Kernel panic - not syncing: sysrq triggered crash
[ 79.804578][ T969] CPU: 6 PID: 969 Comm: bash Not tainted 6.2.0-rc2 #2 e5c88a38d28ce676364b6ebf0...
[ 79.805639][ T969] Hardware name: QEMU Standard PC (Q35 + ICH9, 2009), BIOS 1.15.0-1 04/01/2014
[ 79.806640][ T969] Call Trace:
[ 79.806950][ T969] <TASK>
[ 79.807229][ T969] dump_stack_lvl (lib/dump_stack.c:107 (discriminator 4))
[ 79.807693][ T969] panic (kernel/panic.c:336)
[ 79.808072][ T969] sysrq_handle_crash (drivers/tty/sysrq.c:155)
[ 79.808583][ T969] __handle_sysrq.cold (drivers/tty/sysrq.c:625)
[ 79.809205][ T969] write_sysrq_trigger (drivers/tty/sysrq.c:1163)
[ 79.809668][ T969] proc_reg_write (/arch/x86/include/asm/atomic.h:165
./arch/x86/include/asm/atomic.h
:178 ./include/linux/atomic/atomic-arch-fallback.h:611 ./include/linux/atomic/atomic-instrumented.h:266 f
s/proc/inode.c:211 fs/proc/inode.c:353)
[ 79.810195][ T969] ? rcu_read_lock_any_held (kernel/rcu/update.c:347 kernel/rcu/update.c:340)
[ 79.810750][ T969] vfs_write (fs/read_write.c:582)
[ 79.811157][ T969] ? __do_sys_newfstatat (fs/stat.c:443)
[ 79.811672][ T969] ? lock_is_held_type (kernel/locking/lockdep.c:466 kernel/locking/lockdep.c:5712)
[ 79.812167][ T969] ksys_write (fs/read_write.c:637)
[ 79.812575][ T969] do_syscall_64 (arch/x86/entry/common.c:50 arch/x86/entry/common.c:80)
[ 79.813010][ T969] entry_SYSCALL_64_after_hwframe (arch/x86/entry/entry_64.S:120)
```

Shows the code information of a symbol



**Let's take a deep dive into
how kernel panic actually occurs!**



Analysis kernel panic log

① Skim through Kernel Panic log

② Decode call trace

③ In depth analysis

- Panic trigger command



Analysis kernel panic log

① Skim through Kernel Panic log

② Decode call trace

③ In depth analysis

- Panic trigger command

```
# echo c > /proc/sysrq-trigger
```



Analysis kernel panic log

① Skim through Kernel Panic log

② Decode call trace

③ In depth analysis

- Panic trigger command

```
# echo c > /proc/sysrq-trigger
```

Kernel Proc file

Output redirection




Analysis kernel panic log

① Skim through Kernel Panic log

② Decode call trace

③ In depth analysis

- Panic trigger command



write "c" on /proc/sysrq-trigger



Analysis kernel panic log

① Skim through Kernel Panic log

② Decode call trace

③ In depth analysis

- Panic trigger command

- Analysis with Call Trace



Analysis kernel panic log

① Skim through Kernel Panic log

② Decode call trace

③ In depth analysis

- Panic trigger command

- Analysis with Call Trace

```
[ 79.803379][ T969] sysrq: Trigger a crash
[ 79.803924][ T969] Kernel panic - not syncing: sysrq triggered crash
[ 79.804578][ T969] CPU: 6 PID: 969 Comm: bash Not tainted 6.2.0-rc2 #2 e5c88a38d28ce676364b6ebf0...
[ 79.805639][ T969] Hardware name: QEMU Standard PC (Q35 + ICH9, 2009), BIOS 1.15.0-1 04/01/2014
[ 79.806640][ T969] Call Trace:
[ 79.806950][ T969] <TASK>
[ 79.807229][ T969] dump_stack_lvl (lib/dump_stack.c:107 (discriminator 4))
[ 79.807693][ T969] panic (kernel/panic.c:336)
[ 79.808072][ T969] sysrq_handle_crash (drivers/tty/sysrq.c:155)
[ 79.808583][ T969] __handle_sysrq_cold (drivers/tty/sysrq.c:625)
[ 79.809205][ T969] write_sysrq_trigger (drivers/tty/sysrq.c:1163)
[ 79.809668][ T969] proc_reg_write (./arch/x86/include/asm/atomic.h:165 ./arch/x86/include/asm/atomic.h:178 ./include/linux/atomic/atomic-arch-fallback.h:611 ./include/linux/atomic/atomic-instrumented.h:266 fs/proc/inode.c:211 fs/proc/inode.c:353)
[ 79.810195][ T969] ? rcu_read_lock_any_held (kernel/rcu/update.c:347 kernel/rcu/update.c:340)
[ 79.810750][ T969] vfs_write (fs/read_write.c:582)
[ 79.811157][ T969] ? __do_sys_newfstatat (fs/stat.c:443)
[ 79.811672][ T969] ? lock_is_held_type (kernel/locking/lockdep.c:466 kernel/locking/lockdep.c:5712)
[ 79.812167][ T969] ksys_write (fs/read_write.c:637)
[ 79.812575][ T969] do_syscall_64 (arch/x86/entry/common.c:50 arch/x86/entry/common.c:80)
[ 79.813010][ T969] entry_SYSCALL_64_after_hwframe (arch/x86/entry/entry_64.S:120)
[ 79.813581][ T969] RIP: 0033:0x7f6867fcea37
... SNIP ...
[ 79.820728][ T969] </TASK>
[ 79.821439][ T969] Kernel Offset: disabled
[ 79.821895][ T969] ---[ end Kernel panic - not syncing: sysrq triggered crash ]---
```

Executed
last



Executed
first



Analysis with Call Trace

I will analysis except
question mark in Call Trace

```
[ 79.803379][ T969] sysrq: Trigger a crash
[ 79.803924][ T969] Kernel panic - not syncing: sysrq triggered crash
[ 79.804578][ T969] CPU: 6 PID: 969 Comm: bash Not tainted 6.2.0-rc2 #2 e5c88a38d28ce676364b6ebf0...
[ 79.805639][ T969] Hardware name: QEMU Standard PC (Q35 + ICH9, 2009), BIOS 1.15.0-1 04/01/2014
[ 79.806640][ T969] Call Trace:
[ 79.806950][ T969] <TASK>
[ 79.807229][ T969] dump_stack_lvl (lib/dump_stack.c:107 (discriminator 4))
[ 79.807693][ T969] panic (kernel/panic.c:336)
[ 79.808072][ T969] sysrq_handle_crash (drivers/tty/sysrq.c:155)
[ 79.808583][ T969] __handle_sysrq_cold (drivers/tty/sysrq.c:625)
[ 79.809205][ T969] write_sysrq_trigger (drivers/tty/sysrq.c:1163)
[ 79.809668][ T969] proc_reg_write (./arch/x86/include/asm/atomic.h:165 ./arch/x86/include/asm/atomic.h
:178 ./include/linux/atomic/atomic-arch-fallback.h:611 ./include/linux/atomic/atomic-instrumented.h:266 f
s/proc/inode.c:211 fs/proc/inode.c:353)
[ 79.810195][ T969] ? rcu_read_lock_any_held (kernel/rcu/update.c:347 kernel/rcu/update.c:340)
[ 79.810750][ T969] vfs_write (fs/read_write.c:582)
[ 79.811157][ T969] ? __do_sys_newfstatat (fs/stat.c:443)
[ 79.811672][ T969] ? lock_is_held_type (kernel/locking/lockdep.c:466 kernel/locking/lockdep.c:5712)
[ 79.812167][ T969] ksys_write (fs/read_write.c:637)
[ 79.812575][ T969] do_syscall_64 (arch/x86/entry/common.c:50 arch/x86/entry/common.c:80)
[ 79.813010][ T969] entry_SYSCALL_64_after_hwframe (arch/x86/entry/entry_64.S:120)
[ 79.813581][ T969] RIP: 0033:0x7f6867fcea37
... SNIP ...
[ 79.820728][ T969] </TASK>
[ 79.821439][ T969] Kernel Offset: disabled
[ 79.821895][ T969] ---[ end Kernel panic - not syncing: sysrq triggered crash ]---
```

'?' means that the information about this stack entry is probably not reliable.

<https://stackoverflow.com/a/13117401>



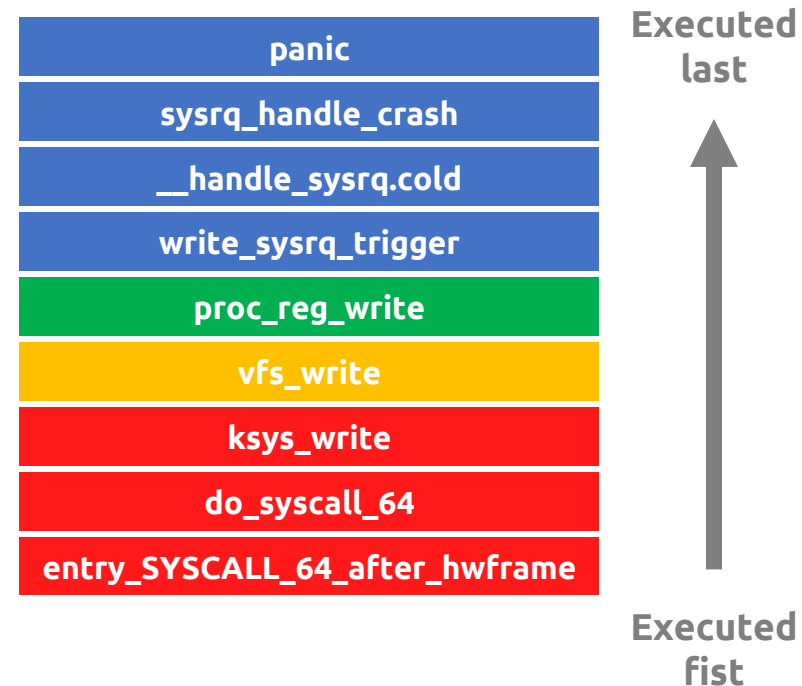
Analysis with Call Trace

```
# echo c > /proc/sysrq-trigger
```

↓
Output redirection

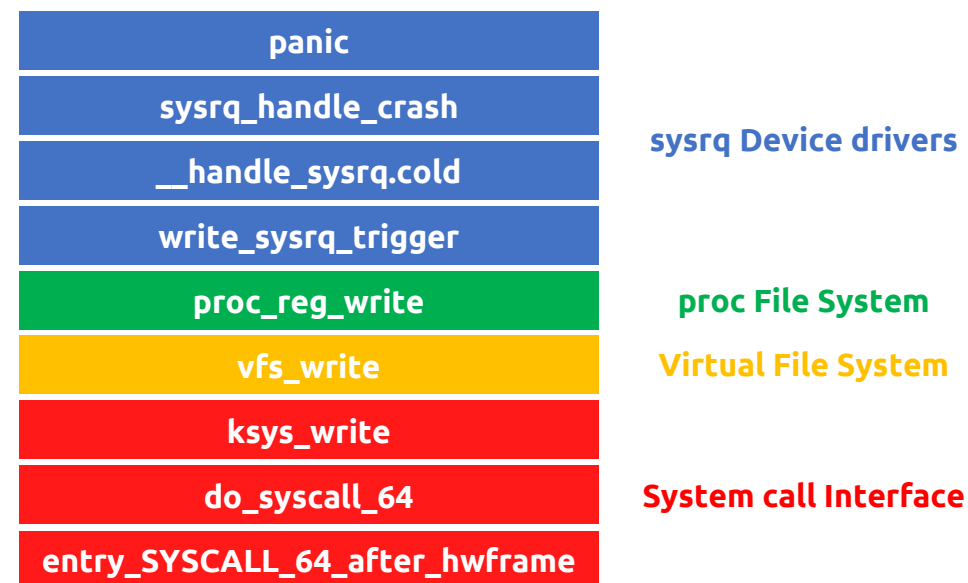
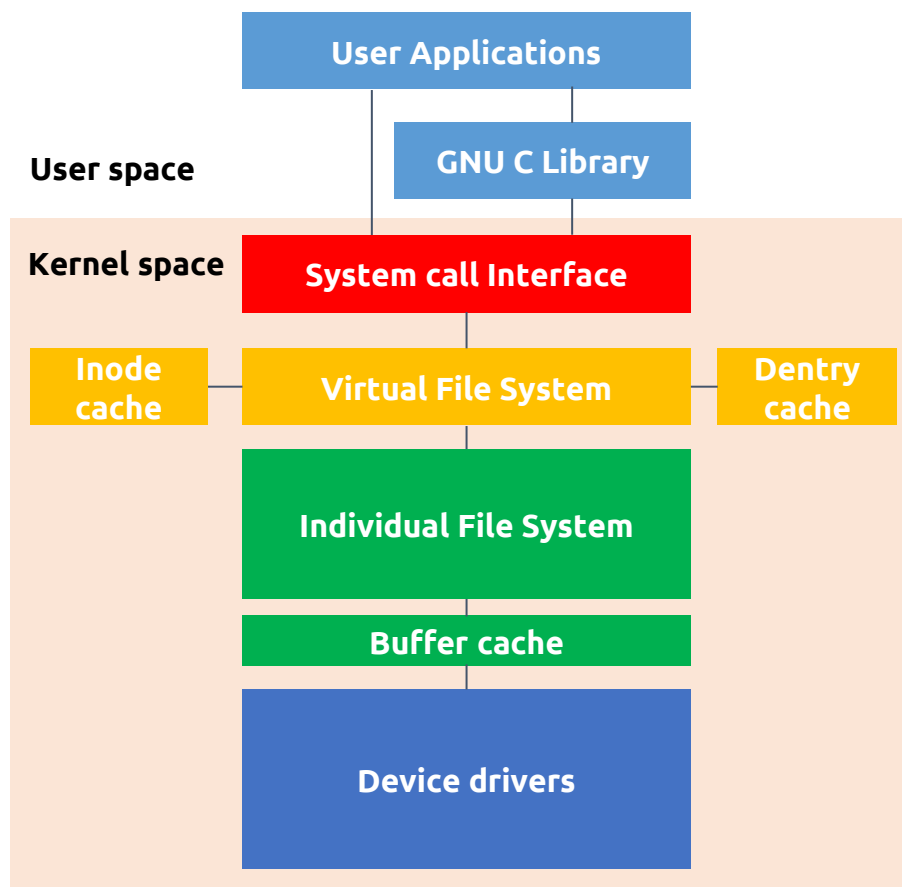
↓
Kernel Proc file

write "c" on /proc/sysrq-trigger





Analysis with Call Trace



<https://developer.ibm.com/tutorials/l-linux-filesystem/>



Analysis with Call Trace

```
[ 79.803379][ T969] sysrq: Trigger a crash
[ 79.803924][ T969] Kernel panic - not syncing: sysrq triggered crash
[ 79.804578][ T969] CPU: 6 PID: 969 Comm: bash Not tainted 6.2.0-rc2 #2 e5c88a38d28ce676364b6ebf0...
[ 79.805639][ T969] Hardware name: QEMU Standard PC (Q35 + ICH9, 2009), BIOS 1.15.0-1 04/01/2014
[ 79.806640][ T969] Call Trace:
[ 79.806950][ T969] <TASK>
[ 79.807229][ T969] dump_stack_lvl (lib/dump_stack.c:107 (discriminator 4))
[ 79.807693][ T969] panic (kernel/panic.c:336)
[ 79.808072][ T969] sysrq_handle_crash (drivers/tty/sysrq.c:155)
[ 79.808583][ T969] __handle_sysrq.cold (drivers/tty/sysrq.c:625)
[ 79.809205][ T969] write_sysrq_trigger (drivers/tty/sysrq.c:1163)
[ 79.809668][ T969] proc_reg_write (/arch/x86/include/asm/atomic.h:165 ./arch/x86/include/asm/atomic.h
:178 ./include/linux/atomic/atomic-arch-fallback.h:611 ./include/linux/atomic/atomic-instrumented.h:266 f
s/proc/inode.c:211 fs/proc/inode.c:353)
[ 79.810195][ T969] ? rcu_read_lock_any_held (kernel/rcu/update.c:347 kernel/rcu/update.c:340)
[ 79.810750][ T969] vfs_write (fs/read_write.c:582)
[ 79.811157][ T969] ? __do_sys_newfstatat (fs/stat.c:443)
[ 79.811672][ T969] ? lock_is_held_type (kernel/locking/lockdep.c:466 kernel/locking/lockdep.c:5712)
[ 79.812167][ T969] ksys_write (fs/read_write.c:637)
[ 79.812575][ T969] do_syscall_64 (arch/x86/entry/common.c:50 arch/x86/entry/common.c:80)
[ 79.813010][ T969] entry_SYSCALL_64_after_hwframe (arch/x86/entry/entry_64.S:120)
[ 79.813581][ T969] RIP: 0033:0x7f6867fcea37
... SNIP ...
[ 79.820728][ T969] </TASK>
[ 79.821439][ T969] Kernel Offset: disabled
[ 79.821895][ T969] ---[ end Kernel panic - not syncing: sysrq triggered crash ]---
```

Let's start using based on 6.2.0-rc2!



entry_SYSCALL_64_after_hwframe

```
106 SYM_INNER_LABEL(entry_SYSCALL_64_after_hwframe,SYM_L_GLOBAL)
107     pushq %rax                /* pt_regs->orig_ax */
108
109     PUSH_AND_CLEAR_REGS rax=$-ENOSYS
110
111     /* IRQs are off. */
112     movq %rsp,%rdi
113     /* Sign extend the lower 32bit as syscall numbers are treated as int */
114     movslq %eax,%rsi
115
116     /* clobbers %rax, make sure it is after saving the syscall nr */
117     IBRS_ENTER
118     UNTRAIN_RET
119
120     call do_syscall_64        /* returns with IRQs disabled */
121
```

arch/x86/entry/entry_64.S

panic	
sysrq_handle_crash	sysrq Device drivers
__handle_sysrq.cold	
write_sysrq_trigger	
proc_reg_write	proc File System
vfs_write	Virtual File System
ksys_write	
do_syscall_64	System call Interface
entry_SYSCALL_64_after_hwframe	



do_syscall_64

```
73 __visible noinstr void do_syscall_64(struct pt_regs *regs, int nr)
74 {
75     add_random_kstack_offset();
76     nr = syscall_enter_from_user_mode(regs, nr);
77
78     instrumentation_begin();
79
80     if (!do_syscall_x64(regs, nr) && !do_syscall_x32(regs, nr) && nr != -1) {
81         /* Invalid system call, but still a system call. */
82         regs->ax = __x64_sys_ni_syscall(regs);
83     }
84
85     instrumentation_end();
86     syscall_exit_to_user_mode(regs);
87 }
```

arch/x86/entry/common.c

panic	
sysrq_handle_crash	sysrq Device drivers
__handle_sysrq.cold	
write_sysrq_trigger	
proc_reg_write	proc File System
vfs_write	Virtual File System
ksys_write	
do_syscall_64	System call Interface
entry_SYSCALL_64_after_hwframe	



do_syscall_64

```
73 __visible noinstr void do_syscall_64(struct pt_regs *regs, int nr)
74 {
75     add_random_kstack_offset();
76     nr = syscall_enter_from_user_mode(regs, nr);
77
78     instrumentation_begin();
79
80     if (!do_syscall_x64(regs, nr) && !do_syscall_x32(regs, nr) && nr != -1) {
81         /* Invalid system call, but still a system call. */
82         regs->ax = __x64_sys_ni_syscall(regs);
83     }
```

```
40 static __always_inline bool do_syscall_x64(struct pt_regs *regs, int nr)
41 {
... SNIP ...
46     unsigned int unr = nr;
47
48     if (likely(unr < NR_syscalls)) {
49         unr = array_index_nospec(unr, NR_syscalls);
50         regs->ax = sys_call_table[unr](regs);
51         return true;
52     }
53     return false;
54 }
```

arch/x86/entry/common.c

panic	
sysrq_handle_crash	sysrq Device drivers
__handle_sysrq.cold	
write_sysrq_trigger	
proc_reg_write	proc File System
vfs_write	Virtual File System
ksys_write	
do_syscall_64	System call Interface
entry_SYSCALL_64_after_hwframe	



do_syscall_64

```
40 static __always_inline bool do_syscall_x64(struct pt_regs *regs, int nr)
41 {
... SNIP ...
50     regs->ax = sys_call_table[unr](regs);
... SNIP ...
```

arch/x86/entry/common.c

```
16 asmlinkage const sys_call_ptr_t sys_call_table[] = {
17 #include <asm/syscalls_64.h>
18 };
```

arch/x86/entry/syscall_64.c

```
1 __SYSCALL(0, sys_read)
2 __SYSCALL(1, sys_write)
3 __SYSCALL(2, sys_open)
4 __SYSCALL(3, sys_close)
5 __SYSCALL(4, sys_newstat)
6 __SYSCALL(5, sys_newfstat)
7 __SYSCALL(6, sys_newlstat)
... SNIP ...
```

arch/x86/include/generated/asm/syscalls_64.h

panic	
sysrq_handle_crash	sysrq Device drivers
__handle_sysrq.cold	
write_sysrq_trigger	
proc_reg_write	proc File System
vfs_write	Virtual File System
ksys_write	
do_syscall_64	System call Interface
entry_SYSCALL_64_after_hwframe	



do_syscall_64

```
40 static __always_inline bool do_syscall_x64(struct pt_regs *regs, int nr)
41 {
... SNIP ...
50     regs->ax = sys_call_table[unr](regs);
... SNIP ...
```

arch/x86/entry/common.c

```
16 asmlinkage const sys_call_ptr_t sys_call_table[] = {
17 #include <asm/syscalls_64.h>
18};
```

arch/x86/entry/syscall_64.c

```
1 __SYSCALL(0, sys_read)
2 __SYSCALL(1, sys_write)
3 __SYSCALL(2, sys_open)
4 __SYSCALL(3, sys_close)
5 __SYSCALL(4, sys_newstat)
6 __SYSCALL(5, sys_newfstat)
7 __SYSCALL(6, sys_newlstat)
... SNIP ...
```

arch/x86/include/generated/asm/syscalls_64.h

```
646 SYSCALL_DEFINE3(write, unsigned int, fd, const char __user *, buf,
647                 size_t, count)
648 {
649     return ksys_write(fd, buf, count);
650 }
```

fs/write.c

panic	
sysrq_handle_crash	sysrq Device drivers
__handle_sysrq.cold	
write_sysrq_trigger	
proc_reg_write	proc File System
vfs_write	Virtual File System
ksys_write	
do_syscall_64	System call Interface
entry_SYSCALL_64_after_hwframe	



do_syscall_64

```
73 __visible noinstr void do_syscall_64(struct pt_regs *regs, int nr)
74 {
75     add_random_kstack_offset();
76     nr = syscall_enter_from_user_mode(regs, nr);
77
78     instrumentation_begin();
79
80     if (!do_syscall_x64(regs, nr) && !do_syscall_x32(regs, nr) && nr != -1) {
81         /* Invalid system call, but still a system call. */
82         regs->ax = __x64_sys_ni_syscall(regs);
83     }
```

```
40 static __always_inline bool do_syscall_x64(struct pt_regs *regs, int nr)
41 {
... SNIP ...
46     unsigned int unr = nr;
47
48     if (likely(unr < NR_syscalls)) {
49         unr = array_index_nospec(unr, NR_syscalls);
50         regs->ax = sys_call_table[unr](regs);
51         return true;
52     }
53     return false;
54 }
```

↑ calls ksys_write

arch/x86/entry/common.c

```
646 SYSCALL_DEFINE3(write, unsigned int, fd, const char __user *, buf,
647                 size_t, count)
648 {
649     return ksys_write(fd, buf, count);
650 }
```

fs/write.c

panic	
sysrq_handle_crash	sysrq Device drivers
__handle_sysrq.cold	
write_sysrq_trigger	
proc_reg_write	proc File System
vfs_write	Virtual File System
ksys_write	
do_syscall_64	System call Interface
entry_SYSCALL_64_after_hwframe	



ksys_write

```
626 ssize_t ksys_write(unsigned int fd, const char __user *buf, size_t count)
627 {
628     struct fd f = fdget_pos(fd);
629     ssize_t ret = -EBADF;
630
631     if (f.file) {
632         loff_t pos, *ppos = file_ppos(f.file);
633         if (ppos) {
634             pos = *ppos;
635             ppos = &pos;
636         }
637         ret = vfs_write(f.file, buf, count, ppos);
638         if (ret >= 0 && ppos)
639             f.file->f_pos = pos;
640         fdput_pos(f);
641     }
642
643     return ret;
644 }
```

fs/read_write.c

panic	
sysrq_handle_crash	sysrq Device drivers
__handle_sysrq.cold	
write_sysrq_trigger	
proc_reg_write	proc File System
vfs_write	Virtual File System
ksys_write	
do_syscall_64	System call Interface
entry_SYSCALL_64_after_hwframe	



vfs_write

```
564 ssize_t vfs_write(struct file *file, const char __user *buf, size_t count, loff_t *pos)
565 {
566     ssize_t ret;
567     ... SNIP ...
575     ret = rw_verify_area(WRITE, file, pos, count);
576     if (ret)
577         return ret;
578     if (count > MAX_RW_COUNT)
579         count = MAX_RW_COUNT;
580     file_start_write(file);
581     if (file->f_op->write)
582         ret = file->f_op->write(file, buf, count, pos);
583     else if (file->f_op->write_iter)
584         ret = new_sync_write(file, buf, count, pos);
585     else
586         ret = -EINVAL;
587     if (ret > 0) {
588         fsnotify_modify(file);
589         add_wchar(current, ret);
590     }
591     inc_syscw(current);
592     file_end_write(file);
593     return ret;
594 }
```

fs/read_write.c

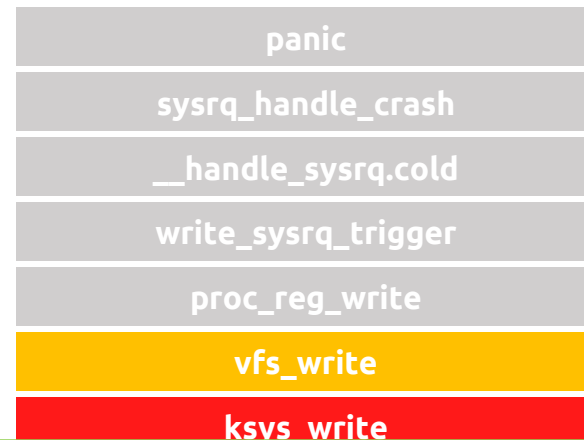
panic	
sysrq_handle_crash	sysrq Device drivers
__handle_sysrq.cold	
write_sysrq_trigger	
proc_reg_write	proc File System
vfs_write	Virtual File System
ksys_write	
do_syscall_64	System call Interface
entry_SYSCALL_64_after_hwframe	



vfs_write

```
564 ssize_t vfs_write(struct file *file, const char __user *buf, size_t count, loff_t *pos)
565 {
566     ssize_t ret;
... SNIP ...
575     ret = rw_verify_area(WRITE, file, pos, count);
576     if (ret)
577         return ret;
578     if (count > MAX_RW_COUNT)
579         count = MAX_RW_COUNT;
580     file_start_write(file);
581     if (file->f_op->write) ↓ calls proc_reg_write
582         ret = file->f_op->write(file, buf, count, pos);
583     else if (file->f_op->write_iter)
584         ret = new_sync_write(file, buf, count, pos);
585     else
586         ret = -EINVAL;
587     if (ret > 0) {
588         fsnotify_modify(file);
589         add_wchar(current, ret);
590     }
591     inc_syscw(current);
592     file_end_write(file);
593     return ret;
594 }
```

fs/read_write.c



sysrq Device drivers

proc File System

Virtual File System

System call Interface

```
578 static const struct file_operations proc_reg_file_ops = {
579     .llseek      = proc_reg_llseek,
580     .read        = proc_reg_read,
581     .write       = proc_reg_write,
582     .poll        = proc_reg_poll,
583     .unlocked_ioctl = proc_reg_unlocked_ioctl,
584     .mmap        = proc_reg_mmap,
585     .get_unmapped_area = proc_reg_get_unmapped_area,
586     .open        = proc_reg_open,
587     .release     = proc_reg_release,
588 };
```

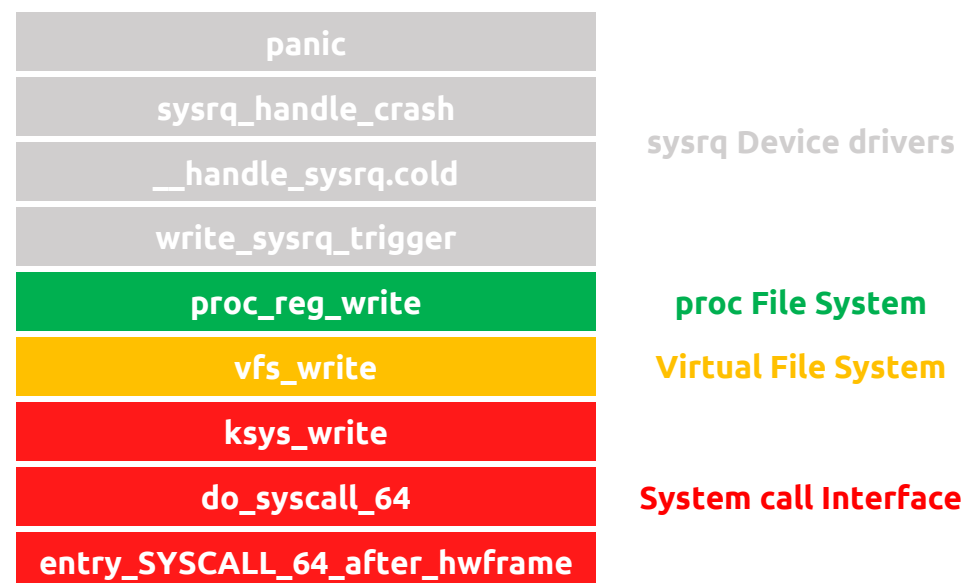
fs/proc/inode.c



proc_reg_write

```
344 static ssize_t proc_reg_write(struct file *file, const char __user *buf, size_t count, loff_t *ppos)
345 {
346     struct proc_dir_entry *pde = PDE(file_inode(file));
347     ssize_t rv = -EIO;
348
349     if (pde_is_permanent(pde)) {
350         return pde_write(pde, file, buf, count, ppos);
351     } else if (use_pde(pde)) {
352         rv = pde_write(pde, file, buf, count, ppos);
353         unuse_pde(pde);
354     }
355     return rv;
356 }
```

fs/proc/inode.c



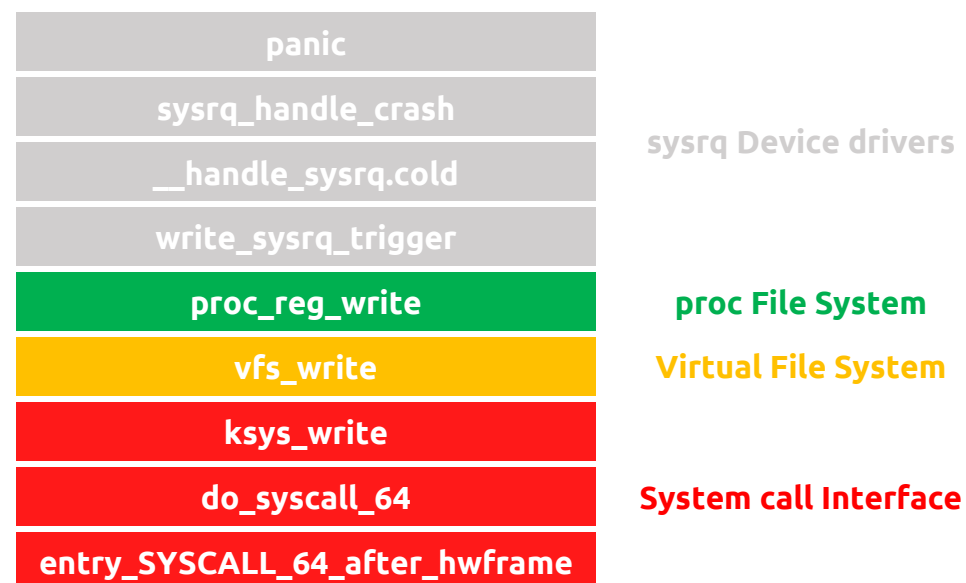


proc_reg_write

```
344 static ssize_t proc_reg_write(struct file *file, const char __user *buf, size_t count, loff_t *ppos)
345 {
346     struct proc_dir_entry *pde = PDE(file_inode(file));
347     ssize_t rv = -EIO;
348
349     if (pde_is_permanent(pde)) {
350         return pde_write(pde, file, buf, count, ppos);
351     } else if (use_pde(pde)) {
352         rv = pde_write(pde, file, buf, count, ppos);
353         unuse_pde(pde);
354     }
355     return rv;
356 }
```

```
334 static ssize_t pde_write(struct proc_dir_entry *pde, struct file *file, const char __user *buf,
size_t count, loff_t *ppos)
335 {
336     typeof_member(struct proc_ops, proc_write) write;
337
338     write = pde->proc_ops->proc_write;
339     if (write)
340         return write(file, buf, count, ppos);
341     return -EIO;
342 }
```

fs/proc/inode.c





proc_reg_write

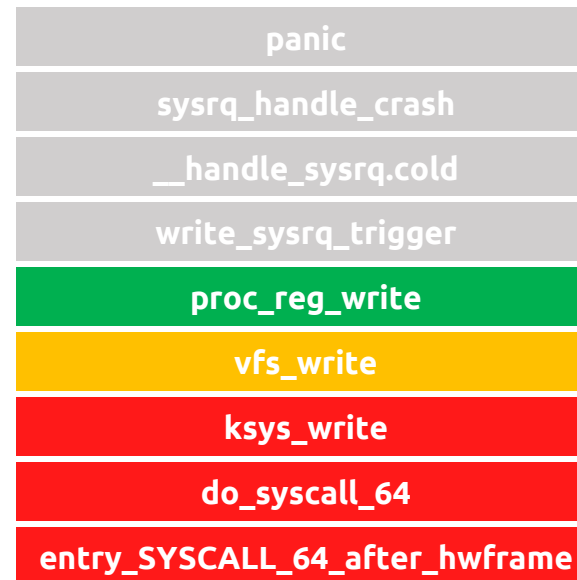
```
344 static ssize_t proc_reg_write(struct file *file, const char __user *buf, size_t count, loff_t *ppos)
345 {
346     struct proc_dir_entry *pde = PDE(file_inode(file));
347     ssize_t rv = -EIO;
348
349     if (pde_is_permanent(pde)) {
350         return pde_write(pde, file, buf, count, ppos);
351     } else if (use_pde(pde)) {
352         rv = pde_write(pde, file, buf, count, ppos);
353         unuse_pde(pde);
354     }
355     return rv;
356 }
```

```
334 static ssize_t pde_write(struct proc_dir_entry *pde, struct file *file, const char __user *buf,
size_t count, loff_t *ppos)
335 {
336     typeof_member(struct proc_ops, proc_write) write;
337
338     write = pde->proc_ops->proc_write; ← calls write_sysrq_trigger
339     if (write)
340         return write(file, buf, count, ppos);
341     return -EIO;
342 }
```

fs/proc/inode.c

```
1169 static const struct proc_ops sysrq_trigger_proc_ops = {
1170     .proc_write      = write_sysrq_trigger,
1171     .proc_lseek      = noop_llseek,
1172 };
```

drivers/tty/sysrq.c



sysrq Device drivers

proc File System

Virtual File System

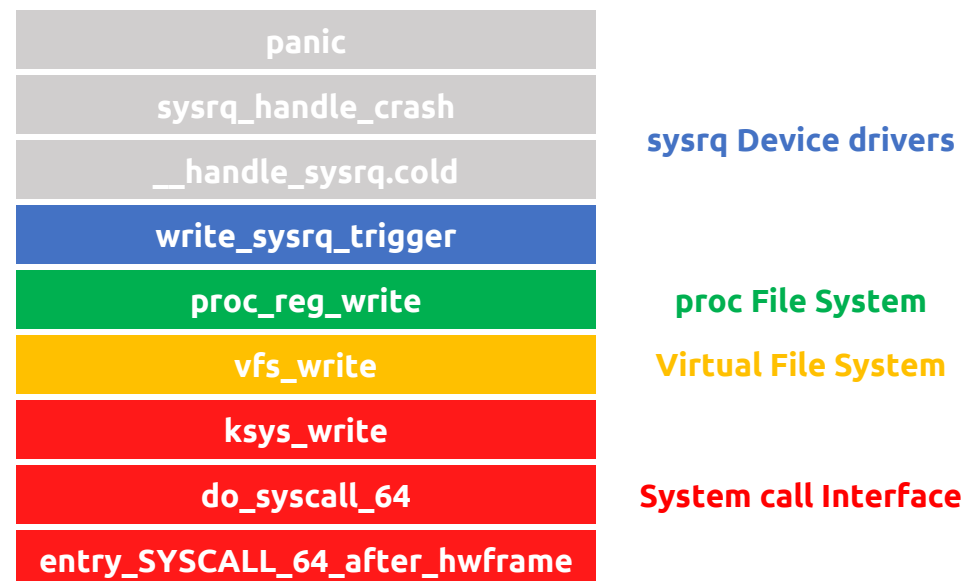
System call Interface



write_sysrq_trigger

```
1155 static ssize_t write_sysrq_trigger(struct file *file, const char __user *buf,  
1156                                     size_t count, loff_t *ppos)  
1157 {  
1158     if (count) {  
1159         char c;  
1160  
1161         if (get_user(c, buf))  
1162             return -EFAULT;  
1163         __handle_sysrq(c, false);  
1164     }  
1165     return count;  
1167 }
```

drivers/tty/sysrq.c





__handle_sysrq

```
572 void __handle_sysrq(int key, bool check_mask)
573 {
574     const struct sysrq_key_op *op_p;
575     int orig_log_level;
576     int orig_suppress_printk;
577     int i;
... SNIP ...
593     op_p = __sysrq_get_key_op(key);
594     if (op_p) {
595         /*
596          * Should we check for enabled operations (/proc/sysrq-trigger
597          * should not) and is the invoked operation enabled?
598          */
599         if (!check_mask || sysrq_on_mask(op_p->enable_mask)) {
600             pr_info("%s\n", op_p->action_msg);
601             console_loglevel = orig_log_level;
602             op_p->handler(key);
603         } else {
604             pr_info("This sysrq operation is disabled.\n");
605             console_loglevel = orig_log_level;
606         }
    }
```

drivers/tty/sysrq.c

panic	
sysrq_handle_crash	
__handle_sysrq.cold	sysrq Device drivers
write_sysrq_trigger	
proc_reg_write	proc File System
vfs_write	Virtual File System
ksys_write	
do_syscall_64	System call Interface
entry_SYSCALL_64_after_hwframe	



__handle_sysrq

```
572 void __handle_sysrq(int key, bool check_mask)
573 {
574     const struct sysrq_key_op *op_p;
575     int orig_log_level;
576     int orig_suppress_printk;
577     int i;
... SNIP ...
593     op_p = __sysrq_get_key_op(key);
594     if (op_p) {
595         /*
596          * Should we check for enabled operations (/proc/sysrq-trigger
597          * should not) and is the invoked operation enabled?
598          */
599         if (!check_mask || sysrq_on_mask(op_p->enable_mask)) {
600             pr_info("%s\n", op_p->action_msg);
601             console_loglevel = orig_log_level;
602             op_p->handler(key);
603         } else {
604             pr_info("This sysrq operation is disabled.\n");
605             console_loglevel = orig_log_level;
606         }
    }
```

drivers/tty/sysrq.c

panic
sysrq_handle_crash
__handle_sysrq.cold
write_sysrq_trigger

sysrq Device drivers

```
552 static const struct sysrq_key_op *__sysrq_get_key_op(int key)
553 {
554     const struct sysrq_key_op *op_p = NULL;
555     int i;
556
557     i = sysrq_key_table_key2index(key);
558     if (i != -1)
559         op_p = sysrq_key_table[i];
560
561     return op_p;
562 }
```

/proc File System

Virtual File System

System call Interface



__handle_sysrq

```
572 void __handle_sysrq(int key, bool check_mask)
573 {
574     const struct sysrq_key_op *op_p;
575     int orig_log_level;
576     int orig_suppress_printk;
577     int i;
... SNIP ...
593     op_p = __sysrq_get_key_op(key);
594     if (op_p) {
595         /*
596          * Should we check for enabled operations (/proc/sysrq-trigger
597          * should not) and is the invoked operation enabled?
598          */
599         if (!check_mask || sysrq_on_mask(op_p->enable_mask)) {
600             pr_info("%s\n", op_p->action_msg);
601             console_loglevel = orig_log_level;
602             op_p->handler(key);
603         } else {
604             pr_info("This sysrq operation is disabled.\n");
605             console_loglevel = orig_log_level;
606         }
    }
```

drivers/tty/sysrq.c

```
552 static const struct sysrq_key_op *__sysrq_get_key_op(int key)
553 {
554     const struct sysrq_key_op *op_p = NULL;
555     int i;
556
557     i = sysrq_key_table_key2index(key);
558     if (i != -1)
559         op_p = sysrq_key_table[i];
560
561     return op_p;
562 }
```

```
456 static const struct sysrq_key_op *sysrq_key_table[62] = {
...
474     &sysrq_crash_op,           /* c */
475     &sysrq_showlocks_op,       /* d */
476     &sysrq_term_op,            /* e */
477     &sysrq_moom_op,            /* f */
478     /* g: May be registered for the kernel debugger */
    }
```

panic
sysrq_handle_crash
__handle_sysrq.cold
write_sysrq_trigger

sysrq Device drivers

/proc File System

Virtual File System

System call Interface

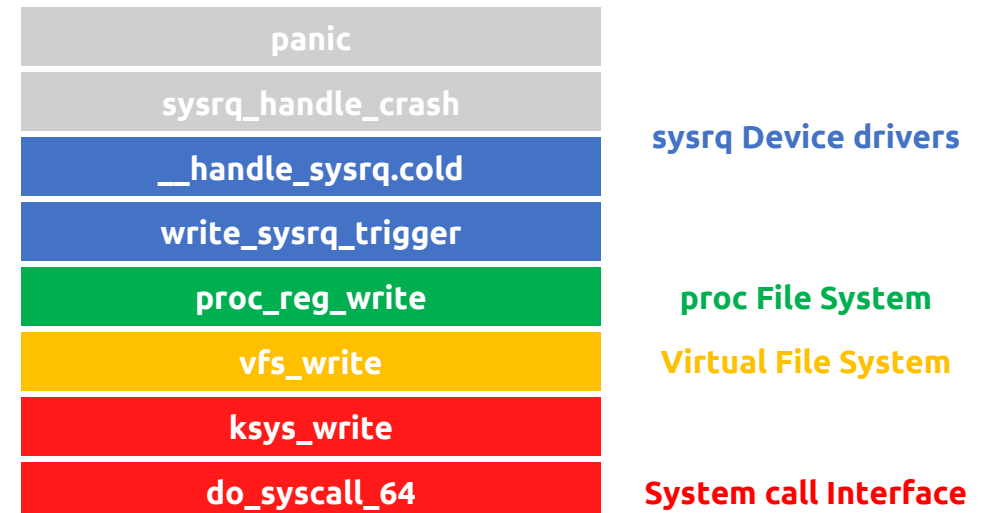


__handle_sysrq

```
572 void __handle_sysrq(int key, bool check_mask)
573 {
574     const struct sysrq_key_op *op_p;
575     int orig_log_level;
576     int orig_suppress_printk;
577     int i;
... SNIP ...
593     op_p = __sysrq_get_key_op(key);
594     if (op_p) {
595         /*
596          * Should we check for enabled operations (/proc/sysrq-trigger
597          * should not) and is the invoked operation enabled?
598          */
599         if (!check_mask || sysrq_on_mask(op_p->enable_mask)) {
600             pr_info("%s\n", op_p->action_msg);
601             console_loglevel = orig_log_level;
602             op_p->handler(key); ← calls sysrq_handle_crash
603         } else {
604             pr_info("This sysrq operation is disabled.\n");
605             console_loglevel = orig_log_level;
606         }
    }
```

drivers/tty/sysrq.c

```
157 static const struct sysrq_key_op sysrq_crash_op = {
158     .handler      = sysrq_handle_crash,
159     .help_msg     = "crash(c)",
160     .action_msg   = "Trigger a crash",
161     .enable_mask  = SYSRQ_ENABLE_DUMP,
162 };
```

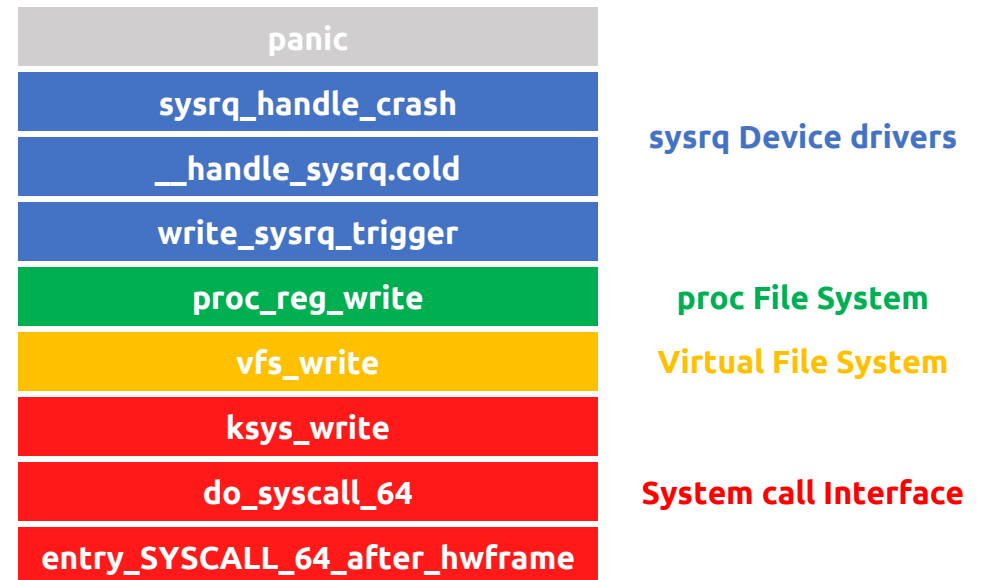




sysrq_handle_crash

```
150 static void sysrq_handle_crash(int key)
151 {
152     /* release the RCU read lock before crashing */
153     rcu_read_unlock();
154
155     panic("sysrq triggered crash\n");
156 }
```

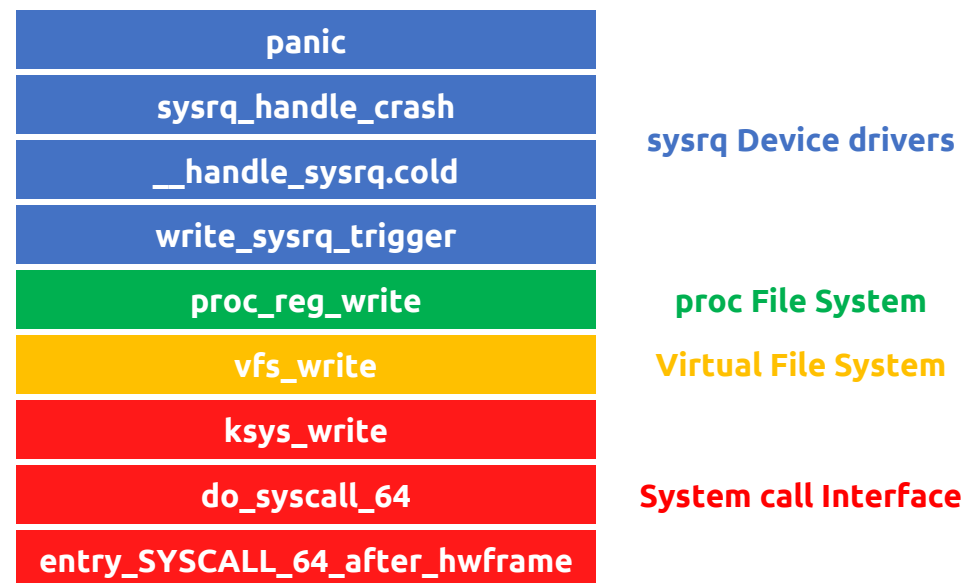
drivers/tty/sysrq.c





```
150 static void sysrq_handle_crash(int key)
151 {
152     /* release the RCU read lock before crashing */
153     rcu_read_unlock();
154
155     panic("sysrq triggered crash\n");
156 }
```

drivers/tty/sysrq.c





Summary

```
# echo c > /proc/sysrq-trigger
```



entry_SYSCALL_64_after_hwframe

do_syscall_64

do_syscall_x64

sys_call_table[unr]()

ksys_write

vfs_write

file->f_op->write()

proc_reg_write

pde_write

pde->proc_ops->proc_write()

write_sysrq_trigger

__handle_sysrq

op_p->handler();

sysrq_handle_crash

Panic ❌



Analysis kernel panic log


① Skim through Kernel Panic log

② Decode call trace

③ In depth analysis

④ Additional resources

- Tracking the syzbot dashboard



<https://syzkaller.appspot.com/upstream>

Information about the bug

when it started happening








when it last happened

how frequently

more crash reports

etc..

syzbot
Linux

 Open [1155]
  Fixed [4262]
  Invalid [9569]
  Kernel Health
  Bug Lifetimes
  Fuzzing
  Crashes

open (1071):

Title	Repro	Count	Last	Reported	Last activity
INFO: task hung in ext4 evict ea inode	syz	2	4d02h	2h08m	2h08m
kernel BUG in set extent bit	C	3	4h58m	4h57m	4h57m
memory leak in reiserfs parse options (2)	C	1	4d05h	5h04m	5h04m
WARNING in devlink free		3	15h59m	8h11m	5h48m
KASAN: use-after-free Read in io_wq_put_and_exit		12	15h48m	16h46m	16h40m
WARNING: locking bug in perf event task sched in (2)		11	5h32m	18h58m	5h48m
WARNING: refcount bug in v9fs vfs lookup	C	1	4d21h	21h43m	21h43m
KASAN: stack-out-of-bounds Read in filemap_get_read_batch	syz	1	4d22h	22h27m	22h27m
WARNING in vma_complete	syz	2	1d01h	1d07h	1d01h
linux-next boot error: WARNING in drm atomic helper set c...		3	5d11h	1d10h	1d10h
kernel BUG in rxrpc_put_call	C	2	1d21h	1d20h	6h07m
WARNING in bpf_xdp_adjust_tail (4)	C	13	9h39m	1d21h	1d21h
WARNING in shmem evict inode		1	5d22h	1d22h	1d22h
BUG: unable to handle kernel paging request in reiserfs_lookup		11	5d07h	2d00h	2d00h
KASAN: use-after-free Read in io_wq_worker_wake	C	1	2d08h	2d08h	1d20h
WARNING in mbind range	C	3	2d01h	2d08h	21h30m
KASAN: slab-out-of-bounds Read in get_block (2)		1	2d14h	2d08h	2d08h
general protection fault in ntfs_lookup		181	3h04m	2d09h	2d09h
general protection fault in reiserfs_get_block		2	5d01h	2d09h	2d09h
general protection fault in udf_fiiter_write_fi		1	7d16h	2d09h	2d09h
KASAN: invalid-free in dbUnmount		1	2d20h	2d09h	2d09h



**With troubleshooting kernel panic,
I could improve my contribution skills!**





Q & A

Juhee Kang(claudiajkang@gmail.com)